

Chat Transcripts

Two message threads — transcribed verbatim with inline images

Source: two screen recordings captured 24 June 2026 (~11:51–11:56). Two separate conversations between Ben (benjamin.pauls) and Xijun Zhang: one in Slack, one in WhatsApp. All wording, spelling and timestamps are reproduced exactly as shown on screen. Times appear as displayed (Slack 12-hour; WhatsApp 24-hour). Shared images are embedded at the point they appear. Resolved calendar dates are added in parentheses on date dividers for reference.

Chat 1 — Slack direct message: Xijun Zhang (VIP) ↔ benjamin.pauls

— Jun 22nd —

Xijun Zhang 1:43 PM

Hey Ben how are you? Good weekend?

Can you share a copy of the Merceros ontology with me?

So for context we need to do an estimate of how long the data foundation timeline would be with Merceros, and right now neither Lorenz or Tom know and advised to understand from Merceros, which Steffen is not a big fan of.

Xijun Zhang 1:46 PM

Steffen also asked for your presence in a session this evening on Copilot front end design. In case you were wondering what the session is about!

↳ Thread: 1 reply — Jun 22nd at 1:56 PM

Ben 1:49 PM

Thank you for the info i will share the ontology with you!

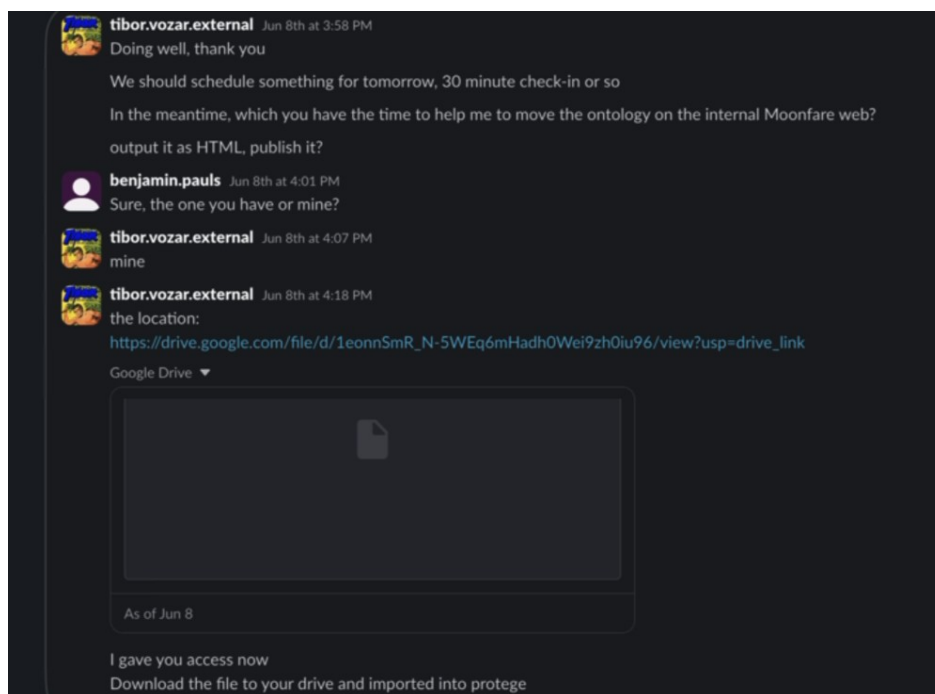
Im good how are you? are you back in the UK?

Ben 1:53 PM

This is what i was given by Tibor as the moonfare ontology

File: MoonfareBiz_v2.ttl — 3.7 MB Turtle

Forwarded into the chat — an earlier Slack exchange (Jun 8th) between tibor.vozar.external and benjamin.pauls:



Forwarded Slack messages (Jun 8th) shared by Ben — Tibor providing the MoonfareBiz_v2.ttl ontology via Google Drive

Transcript of the forwarded block — tibor.vozar.external (Jun 8th, 3:58 PM): “Doing well, thank you” · “We should schedule something for tomorrow, 30 minute check-in or so” · “In the meantime, which you have the time to help me to move the ontology on the internal Moonfare web?” · “output it as HTML, publish it?” | benjamin.pauls (4:01 PM): “Sure, the one you have or mine?” | tibor.vozar.external (4:07 PM): “mine” | tibor.vozar.external (4:18 PM): “the location: https://drive.google.com/file/d/1eonnSmR_N-5WEq6mHadh0Wei9zh0iu96/view?usp=drive_link” [Google Drive: MoonfareBiz_v2.ttl] · “I gave you access now” · “Download the file to your drive and imported into protege” · “Let me know if it works (it should)”.

Xijun Zhang 1:54 PM

Also heard you mapped all the processes of Investment over the weekend?!

↳ Thread: 1 reply — Jun 22nd at 2:00 PM

Xijun Zhang 1:54 PM

You are a machine..

Ben 1:54 PM

Hahaha im not sure how detailed it is or correct used confluence and my own judgement for a lot of it.

Xijun Zhang 1:55 PM

I am back in the UK might not travel so much until Moonrise, are you joining?

Ben 1:55 PM

Dont think i will, leaving Berlin the week before, i am finished with this city.

Reaction: 🍻 (1)

Xijun Zhang 1:56 PM

I'll send you my latest next steps on data on Whatsapp

Ben 1:57 PM

Sounds good!

— Yesterday (Jun 23rd) —

Xijun Zhang 4:57 PM

Hey Ben, how are you?

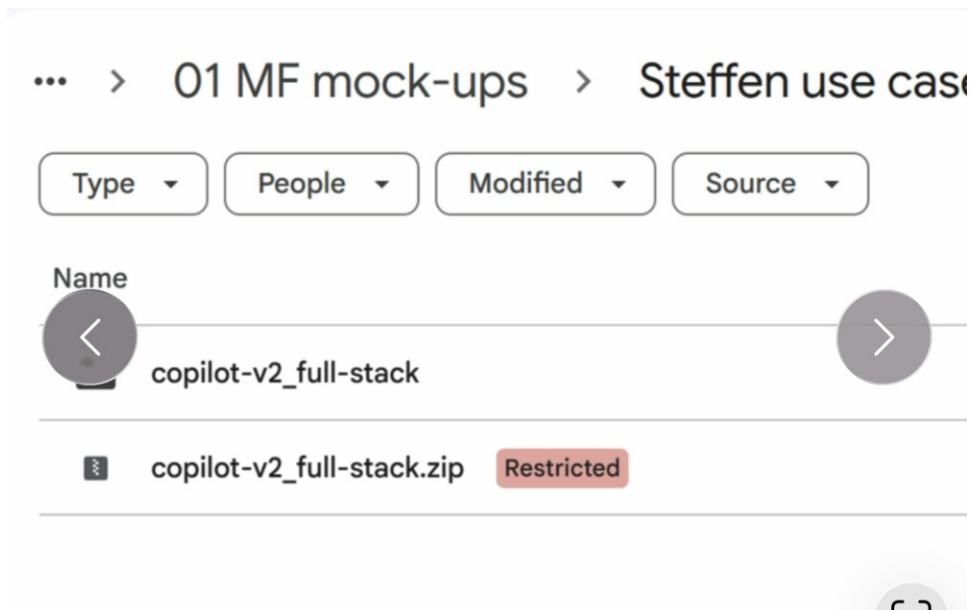
Ben 4:57 PM

Im great how are you?

Xijun Zhang 4:59 PM

Merceros is looking for your UI, would you mind sharing that?

ah i see there is something there



Screenshot shared by Xijun — file browser: "01 MF mock-ups > Steffen use case...", folder "copilot-v2_full-stack", file "copilot-v2_full-stack.zip" (Restricted)

Ben 5:00 PM

Yeah sure ill host it on a cloudflare link free and send it to them.

Xijun Zhang 5:00 PM

no need

just to check this covers everything?

Ben 5:00 PM

Synthetic data based on real data we got from them, frontend UI and logic

Xijun Zhang 5:01 PM

clear thanks!

Ben 5:01 PM

No worries

Xijun Zhang 5:02 PM

Also do you know if Carola is out?

Ben 5:02 PM

I dont but i think she is

Xijun Zhang 5:03 PM

ah ok let me set up the session Steffen requested then

Ben 5:05 PM

I thought Steffen said something about her being out yesterday in the meeting

Xijun Zhang 5:05 PM

ah, do you know his availability? Shall I just set up those meetings direct? What's his preference

Ben 5:06 PM

Carola should be in
just asked him

Xijun Zhang 5:15 PM

☐ I just sent a reminder to Carola for the twice a week review session with Merceros

Ben 5:21 PM

Great!

Xijun Zhang 6:16 PM




Image shared by Xijun (6:16 PM) — "MoonfareBiz Ontology — Logical Modules" (24 modules / ~280 classes, grouped by domain)

Chat 2 — WhatsApp: Ben ↔ Xijun Zhang

In WhatsApp, Ben is the green/right-hand sender; Xijun Zhang is the white/left-hand sender.

— Fri, May 15 —

Messages and calls are end-to-end encrypted. Only people in this chat can read, listen to, or share them. [Learn more](#)
Xijun is a contact.

Xijun Zhang 13:58

□ Xijun here

Ben 14:51

Hey!

— Sat, May 16 —

Xijun Zhang (approx. 11:03)

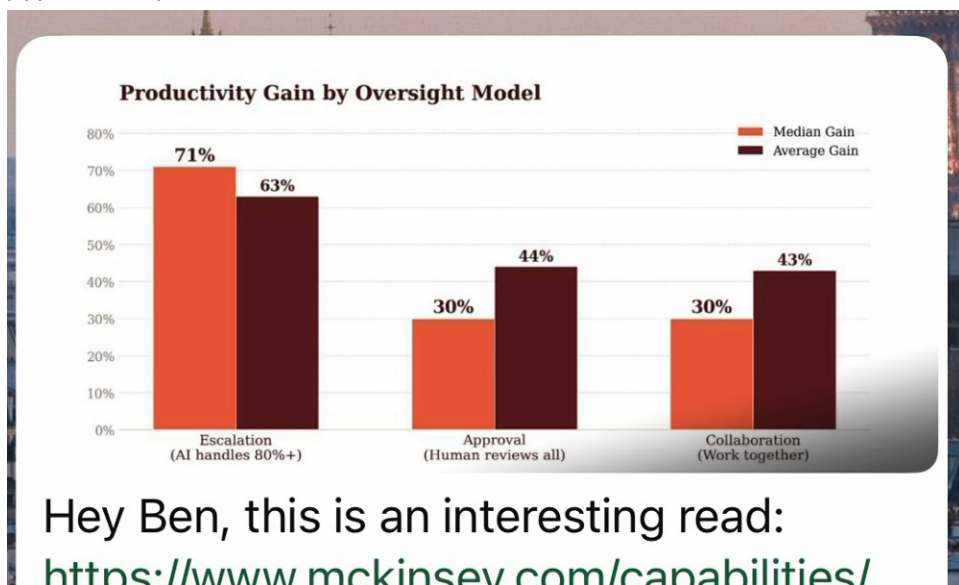


Image shared by Xijun — "Productivity Gain by Oversight Model" (Median vs Average Gain): Escalation 71%/63%, Approval 30%/44%, Collaboration 30%/43%

Xijun Zhang

Hey Ben, this is an interesting read:

<https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-ai-table-stakes-to-ai-advantage-building-competitive-moats>

Xijun Zhang 11:03

I think in answer to your question yesterday what do we really need, is Speed. This is where we need to reduce the level of human interference. The latest Stanford report (https://digitaleconomy.stanford.edu/app/uploads/2026/03/EnterpriseAIPlaybook_PereiraGraylin-Brynjolfsson.pdf) from April already shows fully Agentic is better than Human in the loop (HITL) by a 30% margin. We are a regulated business for sure, we should insist on reviewing sufficiently so they are trained to certain quality before we fully handover, but we should not insist on HITL everywhere, especially where law doesn't reach.

— Mon, May 18 —

Ben 02:04

I couldn't agree more the goal should be regulation yes but the minimum necessary makes total sense 24 hr output at top performance vs realistic 4-5 hours with human error. For a fraction of the cost. Not sure if the EU would allow the training on customer and fund data ☹ in McKinley's report.

I wonder if anonymizing the data gets around that issue. To your point on no one having ownership and fear of decision making, I know according to the ai guideline doc Claude shouldn't make the decision but it could help SWOT it, doesn't solve ownership but may make owning the outcome "less risky"

To your point on engineering and the "external team" McKinsey "Agentic AI is already showing signs of significantly accelerating every stage of the development and deployment cycle, moving from two-week to 24-hour development cycles (or sprints)."

Xijun Zhang 08:44

Good points! On anonymisation, we can use local llm to pseudo anonymise data. TBH, automation is a huge opportunity already, and much of our processes need AI for coordination rather than inferencing to start with, we can first train AI to route based on task than producing better results. This is already better than the traditional automation.

AI native: Speed comes from empowering capable people with delegation of authority to front line. A few challenges here: 1. we need people who knows how to play with agents, 2. we need to change our decision framework, Legal/compliance needs to publish clear decision boundary guidelines on how to empower people to make decisions rather than stopping people every time and saying materials need a review.

Ben 14:14

Have you heard of Odin servers they're being considered instead of MacBook studios as the inference for local.

Ben 17:20

Just had the call with the external providers. Still have no clue what they've been doing. They aren't even connected to the database

Reaction: 🤔

Xijun Zhang 17:23

I was not impressed when I talked to them, and hey you found out first session, it took me 3 sessions to find out

Ben 17:25

They sound impressive because they speak the language to an extent it's just no content at all I have no idea what they've been doing they showed me a chat function to read in files that they curated. I've never seen something made less impressive honestly a skill

Xijun Zhang 17:26

well, it's how you get sold things
not what it actually does but what dream looks like

Ben 17:26

Amazing salesman this guy should get an award

Reaction: 🙄

— Thu, Jun 11 —

Ben 20:53

Just out of curiosity what would we talk about? I'm sure there are more qualified people than me haha.

Xijun Zhang 20:54

you mean at the roadshow?

Ben 20:55

Yes

Xijun Zhang 20:58

It's 8 mins for each person to show what they were able to automate. And it's for those who learnt AI based on self interest. So far the level we are talking about are watermarking tool, financial reporting building agent, process flow mapping skill, support desk bot

Xijun Zhang 20:59

reason why I wanted you to do it is to showcase the art of possible and give the company a bit of pressure

Ben 21:00

Ok that's tamer than what I thought haha. I'm just thinking of the politics of it all haha

Xijun Zhang 21:01

lol what did you think it was?

Xijun Zhang 21:02

but as momentum builds we do need to spice things up a little

Xijun Zhang 21:02

now is not the right time yet

Ben 21:03

To be honest I had no expectations so I thought it would be some people just asking any and all questions

Ben 21:03

↳ replying to Xijun Zhang: "but as momentum builds we do need to spice things up a little"

How would you want to spice it up?

Xijun Zhang 21:04

up the reward + make it known

Ben 21:05

Good idea there are some monitors on both floors in office where you could have a leaderboard

Xijun Zhang 21:07

yeah, I think we might do it after the summer party to keep up the momentum

— Fri, Jun 12 —

Ben 09:58

I like the initiative haha

Xijun Zhang 10:00

do you reckon 3k would be a good reward each quarter? Top place getting 1500, then 1000, then 500? Provided the value generated far exceeds this

Ben 10:17

Yes 100%

— Monday (Jun 22) —

Xijun Zhang 13:59

Hello Sir!

So what came out of last Friday discussion was a few really odd things:

1. we need to rely on 3rd parties to understand our data
2. I found out today our backend team never worked with Merceros
3. we have minimal documentation on data and the data engineer doesn't know 60-80% of what fields mean and how they were created

I will work with Agustin, who has been here a few years and try to get a first view of how long it would take.

Xijun Zhang 14:02

Lukas agrees this is super fishy, not sure what you think. And we just signed another deal with them worth 900k for 3 months of work, I would be laughing my head off inside if I were them.

Ben 14:02

You are kidding me!!! You are messing with me.

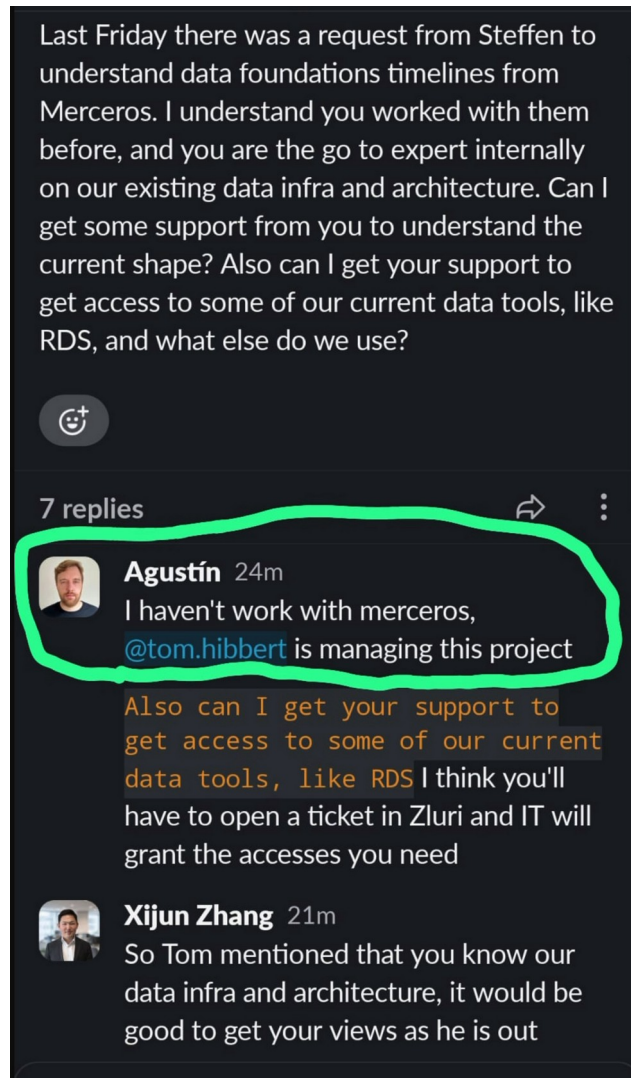
Xijun Zhang 14:02

No

Xijun Zhang 14:02

this is what happened and hope you can explain it in a way that makes sense to Steffen

Xijun Zhang 14:03



Screenshot shared by Xijun — Slack thread (request to Agustín re: data foundations / RDS access). Agustín: "I haven't work with merceros, @tom.hibbert is managing this project ... open a ticket in Zluri and IT will grant the accesses you need." Xijun: "So Tom mentioned that you know our data infra and architecture..."

Xijun Zhang 14:04

If you check this guy leads our backend dev

Ben 14:04

How have they been able to work for the past 6 months it seems like then there was no oversight over anything.

Ben 14:05

BTW if you check the ontology and ask Claude he'll verify that 70% of the ontology is from essentially public ontology sources

Xijun Zhang 14:06

Yeah exactly I don't know tbh Steffen's view on this, I don't want to second guess anything, so I am just presenting the facts here

Xijun Zhang 14:08

Please keep me out of this if you can, as you know I am reporting to Lorenz, but I feel Steffen has 150% of rights to know this

Xijun Zhang 14:09

Also implications here are that we need to find a solution too, and not expose just issues

Ben 14:09

| ↳ replying to Xijun Zhang: "Yeah exactly I don't know tbh Steffen's view on this..."

He told me that he doesn't understand the technical aspects of what is required and therefore can't judge work vs cost of the project. I think that's why also not relying on merceros to tell the timeline and cost of this was brought up. Also from what I understand this is Lorenz's project.

Ben 14:10

| ↳ replying to Xijun Zhang: "Please keep me out of this if you can..."

I'll keep you out of it. I would still of course value your opinion off the record!

Ben 14:10

| ↳ replying to Xijun Zhang: "Also implications here are that we need to find a solution too..."

Of course the past is the past

Xijun Zhang 14:13

Yeah tbh, there is something going on in the background that I am not aware of, and Lorenz thinks Tom being one of the best CTOs he's met. Feels strange but ill stop there

Xijun Zhang 14:13

I feel like opening a pandora box

Ben 14:14

Tom definitely has a lot on his plate. Separate from the whole data part

Xijun Zhang 14:14

TBH for any contracts back in my consulting days, there are clear deliverables for each project so far they haven't got any and we are 3rd week in

Xijun Zhang 14:15

| ↳ replying to Ben: "Tom definitely has a lot on his plate. Separate from the whole data part"

yeah I am staying open minded here too, and seeing where I can help

Ben 14:16

| ↳ replying to Xijun Zhang: "TBH for any contracts back in my consulting days..."

I thought they've been here since February

Xijun Zhang 14:17

so for the first 3 months, they delivered the ontology and a front-end demo

Xijun Zhang 14:18

the way it was done was not via analysing data (I may be wrong don't quote me as they were only recently asking for access) but via a series of expert interviews

Ben 14:18

How was the demo?

Xijun Zhang 14:18

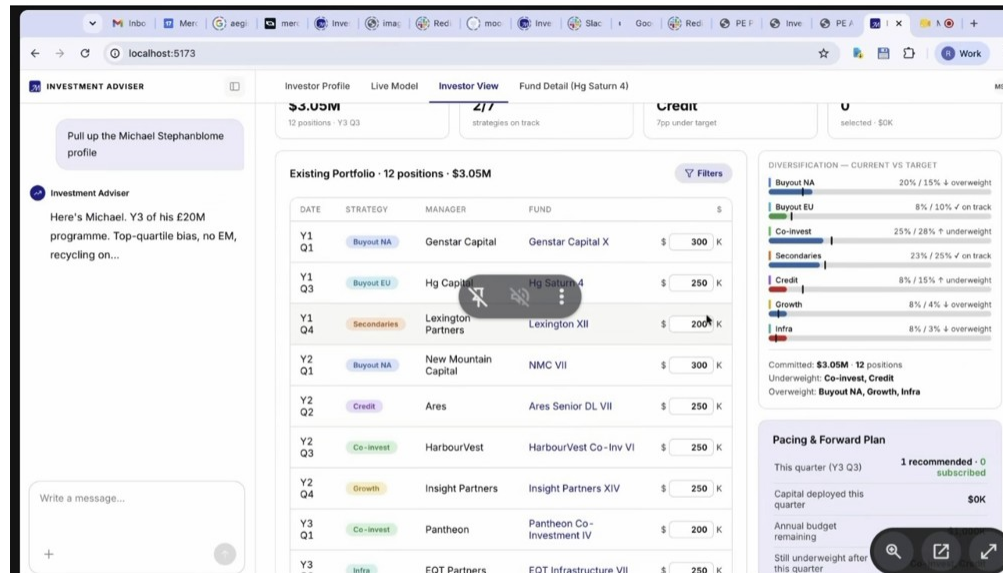
the demo to the board or the one last friday??

Ben 14:19

The one they delivered was the one for Friday right?

Xijun Zhang 14:19

One last Friday was like this. Andy was able to make something that looks like investor profile in 1-2 weeks



Screenshot shared by Xijun (14:19) — "Investment Adviser" portfolio dashboard: Existing Portfolio – 12 positions – \$3.05M

Xijun Zhang 14:20

I mean Andy Todd in Sales, and he is not an expert

Ben 14:20

You are kidding me

Ben 14:20

Claude design?

Xijun Zhang 14:21

this is what andy made

File: mockup-v1.html — 48 KB html



The mockup-v1.html opened — Moonfare “Portfolio Review — Acme Family Office”, prepared by Andy Todd (As of 31 Mar 2026). Total Portfolio Value \$8.42m; Net IRR 18.4%, Net TVPI 1.34x, DPI 0.12x

Ben 14:21



Xijun Zhang 14:21

yeah above was what Merceros presented

Xijun Zhang 14:21

I'm like...

Xijun Zhang 14:22

Now even weirder is that the conversation has been shaped so we need to decide on what backend data we clean based on choice of front end use cases..

Ben 14:22

I thought because the ontology was so bad that they were mainly working on the data integration...

Xijun Zhang 14:23

like we are not making a prd, but we are driving engineering work with front end design...

Xijun Zhang 14:24

the whole thing is a bit messed up to be honest, and Lorenz is shaping this, I don't know what his skin in the game is, but he and Tom seem to have decided things..

Xijun Zhang 14:25
↳ replying to Ben: “I thought because the ontology was so bad that they were mainly working on the data integration...”
haha, I would have thought the same, I don't think there is any cleaning of data yet

Ben 14:25
At least Going forward will there be milestones accountability steps, deadlines, deliverables?

Ben 14:26
↳ replying to Xijun Zhang: “haha, I would have thought the same...”
But what was done then over the past 5 months

Xijun Zhang 14:27
this
File: *AI_CoPilot_Questions (2).pdf* — 2 MB pdf

Q2 — Concrete Outputs Delivered to Date	
Phase 1 (Jan–Apr 2026) — Activities vs. actual deliverables across both tracks	
Track 1 — Data Foundations	Track 2 — AI Co-Pilot Prototype
Formal Ontology for co-pilot (5 Domains) Machine-readable definitions, relationships and rules across Fund, Legal Entity, Investor, Transaction, Portfolio — consumable directly by the AI agent. Done	8 Agent Skills Built Portfolio extraction, investor profiling, peer benchmark construction, cashflow forecasting, gap analysis, full guardrails. Done
Co-pilot Ontology → Production Mapping Field-by-field feasibility validated against real Moonfare production data. Gaps documented. JSON representation of Fund domain built for Co-Pilot. Done	/plan + /ask Interface Structured 15–30 min portfolio planning mode + full guardrails. Scoped to private markets only. Powered by synthetic data. Done
Field-by-Field Data Mapping Vertical slice: Fund domain first (in progress). For each ontology field: do we have it? where? what quality? what breaks if we don't? In Progress	3 Investor Demos + Board Deck Synthetic data for 3 investor personas (incl. 55yo L. Full journey end-to-end. Board deck delivered April Done
Co-pilot Ontology Validation - First Pass Confirmed feasibility of the semantic layer against production systems. Data gaps identified and logged for the pipeline build phase. This needs further validation as is an on-going topic but first pass has been done. Done	

PDF preview shared by Xijun (14:27) — “Q2 — Concrete Outputs Delivered to Date” (Phase 1, Jan–Apr 2026). Track 1 Data Foundations & Track 2 AI Co-Pilot Prototype with Done / In Progress status

Xijun Zhang 14:27
i.e. ontology + copilot

Ben 14:33
I don't understand...

Ben 14:33
↳ replying to Xijun Zhang: “One last Friday was like this. Andy was able to make something that looks like investor profile in 1-2 weeks”
This

Ben 14:33
+ the ontology I sent you

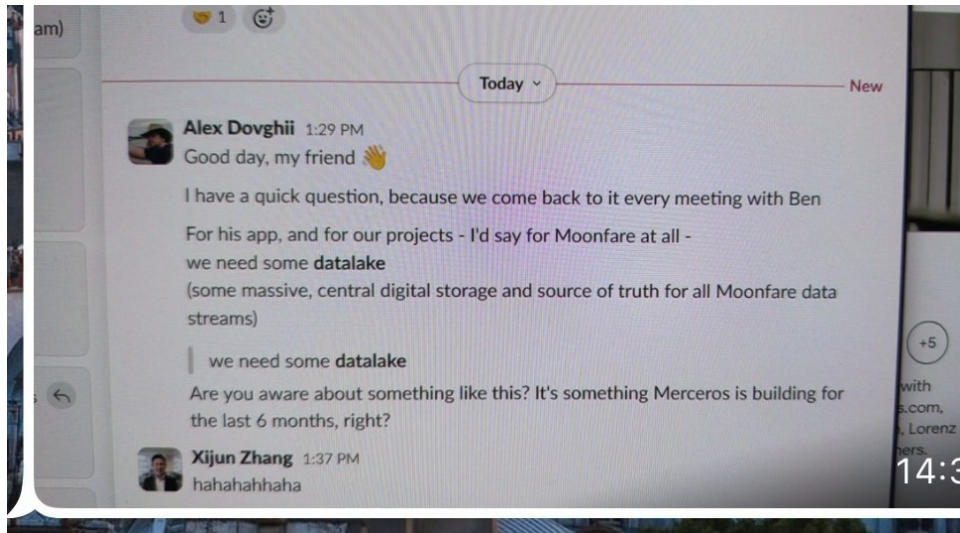
Ben 14:33
↳ replying to Ben: “At least Going forward will there be milestones accountability steps, deadlines, deliverables?”
The dirty politics is that I am now being pushed onto the stage appearing as “owner” of this project. And realistically Lorenz and Tom are shaping this work and directing Merceros in the background. I'll do what I can to uncover what Merceros is doing, and steer the conversation, but what is going on is above me

Ben 14:34
307k?

Xijun Zhang 14:35
And Lorenz is watching me every step

Xijun Zhang 14:35
So it's a bit tricky tbh

Xijun Zhang 14:37



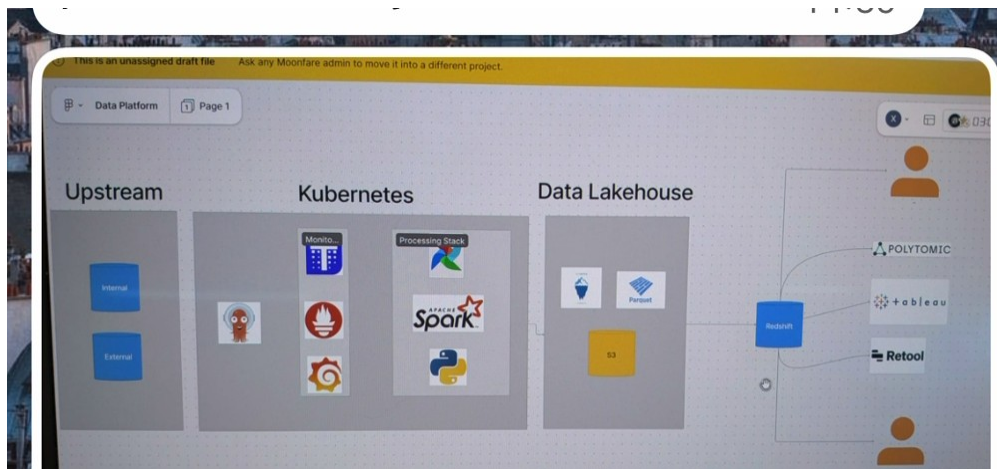
Screenshot shared by Xijun (14:37) — Slack message from Alex Dovghii asking about a “datalake” (central digital storage / source of truth for all Moonfare data) and whether Merceros is building it

Ben 14:38
🤔🤔 I love the enthusiasm

Xijun Zhang 14:38
I'll check with Agustin

Xijun Zhang 14:39
And see our datalake, I think we have one, good now I am speaking with the person that may know

Xijun Zhang 14:40



Screenshot shared by Xijun (14:40) — infrastructure diagram: “Upstream / Kubernetes / Data Lakehouse” (Spark, Python and other tooling)

Xijun Zhang 14:41
RDS is where we aggregate

Ben 14:42

So if I got this right we are then essentially designing the front end plus the functioning of how it interacts and we do that weekly then whilst that is going on Merceros provides the data for that Copilot use case that we are designing the front end for.
Eventually we'll have all data for the copilot and the front end.

Ben 14:42

Do you know if the investments data already lives in the lake house so we could access via RDS?

Xijun Zhang 14:42

eventually we'll have the data for the front end based on the design

Xijun Zhang 14:43

and use case yes

Xijun Zhang 14:43

| ↳ replying to Ben: "Do you know if the investments data already lives in the lake house so we could access via RDS?"
no because gdrive is not connected

Xijun Zhang 14:43

nor are the transcripts

Ben 14:44

Got it so no functionality for that yet.

Ben 14:44

Let me know what Agustin says in terms of how long this should take merceros

Xijun Zhang 14:45

A datalake is at the foundation level, the raw repository we talked about that is simply versioned with some basic Meta data like sources

Xijun Zhang 14:45

we don't have that and we force aggregate at the operational layer

Xijun Zhang 14:46

which is the rds

Xijun Zhang 14:46

the way I see it

Ben 14:46

Got it do you know timeline for that?

Xijun Zhang 14:47

Ben it's a black box, according to Steffen's own words

Xijun Zhang 14:48

I was looking after the enterprise strategy, AI tooling, upskilling and process documentation so far.
But it seems data is the gravest issue

Ben 14:49

| ↳ replying to Xijun Zhang: "Hello Sir! So what came out of last Friday discussion was a few really odd things:"
Last two weeks I found these out

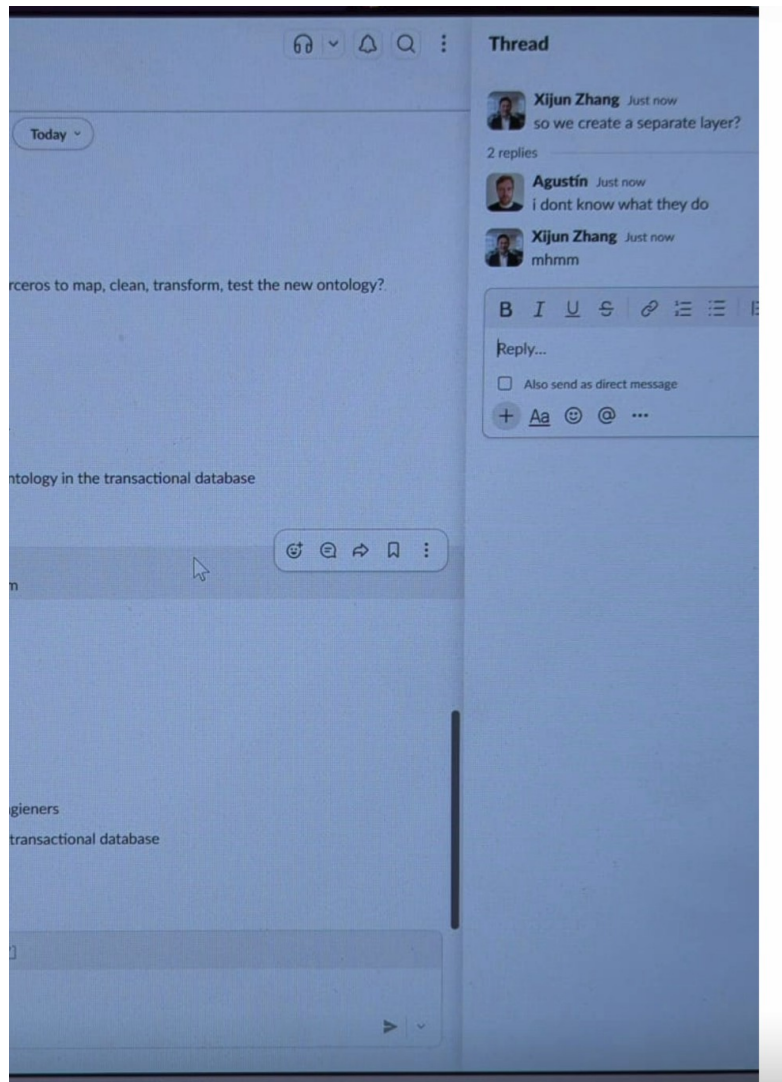
Ben 14:50

Sounds like as you and steffen said Black box and Pandora's box I'm pretty sure the black box is inside of Pandora's box but it seems like Pandora's box is also a black box

Xijun Zhang 14:51

definitely the latter, well put

Xijun Zhang 15:56



Screenshot shared by Xijun (15:56) — Slack thread with Agustín: "the ontology is done and handed over to engieners, then engineers would need to refactor the transactional database ... in my opinion, that is a huge effort." Side thread: "so we create a separate layer?" — Agustín: "i dont know what they do" — Xijun: "mhmm"

Ben 16:02
Lovely

— Yesterday (Jun 23) —

Xijun Zhang 12:49

I saw Steffen's questions for Tibor, good questions!

Ben 13:33

What do you think the feedback will be?

Xijun Zhang 14:16

From my read of Tibor, it will be what they discussed with Tom and they will deliver to the SoW. Then they will say that they can work on what we want them to work on.

The core issue is all their work is outside our system atm, they don't have much outside the ontology, copilot demo. They recently started working on the architecture of their analytical platform, but that's outside our system too.

I am speaking with Agustin later today, but Alex tells me it's 3-5 people's work for 3-4 months to do our entire data layer (please don't quote him as Agustin is the real expert). Merceros have 5 people,

and for 3 months they are delivering a static copilot database that lives outside our system. We need to work with Tom to regroup expectations and support model, it is so strange that neither Agustin or Farooq had much interaction with them.

Xijun Zhang 14:17

As you saw with investment, we have data, but it's not that much, our business model as MSB put it is not difficult

Ben 14:21

I mean if you look at other companies where this has been done even the large hyperscalars take 3 months and comparable fintech companies less. At most 3 months with a team of 3 people

Xijun Zhang 14:21

exactly

Xijun Zhang 14:22

This is Tom's remit and we need him to commit

Xijun Zhang 15:18

Current redshift migration took 4 months with 2 people fyi

Xijun Zhang 15:18

Data I just got from engineers

Ben 15:43

Let's hope some accounting comes from the questions

Xijun Zhang 15:46 (Edited)

This is where I'd best not play any role so I can play a long term more supportive role and actually push the projects forward for the company

Xijun Zhang 15:53

forward for the company

Edited 15:46



Image shared by Xijun (15:53) — an iceberg captioned "Accounting" (tip-of-the-iceberg meme)

Ben 15:54

Like an adversarial LLM session

Xijun Zhang 15:55

I mean after seeing your mock up I don't know why we are paying them

Xijun Zhang 15:56

I can only suspect, but then I am not in the position to

Xijun Zhang 15:56

Your version is not the adversarial anymore, it's a challenge for them to beat

Ben 15:57

Hahaha I appreciate it!

Ben 15:57

I liked Michael's comments yesterday about lowering expectations massively

Reaction: ❤️

Xijun Zhang 15:59

Let's do that

Ben 16:04

Honestly I just don't know where to set the expectations anymore because I would think that the expectation should be at getting the data together but the ontology doesn't even work for the Investment data, I tested it and Phillip kept questioning why my extractor doesn't extract any of the info he needs, built on their ontology. So I don't know where to set my expectations because if it can't even do that then I would hope they spent 99% of their time on the other data but since they haven't even talked to the people from the data team nor do they seem to have access to the data I don't know how that's possible.

Xijun Zhang 18:16

Yeah I get you, we just had an hour call with Merceros, I think they are now clear what they need to deliver (all the data foundation across the 5 highlighted domains in a live system)

Xijun Zhang 18:17

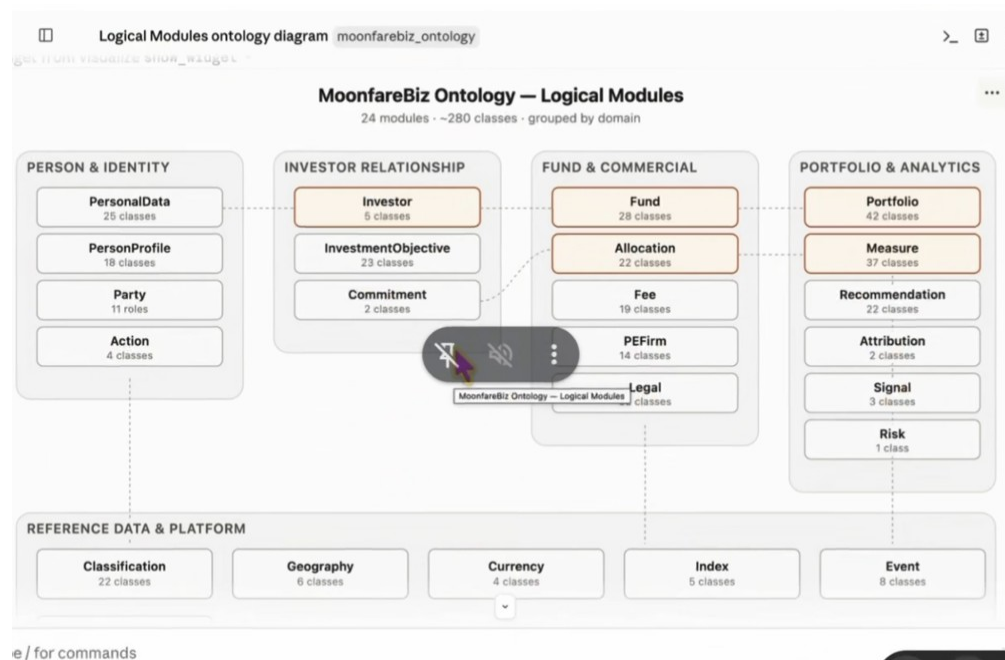


Image shared by Xijun (18:17) — "MoonfareBiz Ontology — Logical Modules" (the same logical-modules ontology diagram)

Xijun Zhang 18:17

All by September

Ben 18:19

Anything else they mentioned?

Ben 18:19

Because 2 of those are essentially done through the investment team data project which is done

Xijun Zhang 18:30

Yes they said they will get it to a high bar (TBD what that means) and get all of the data live so we can use it

Xijun Zhang 18:30

it will be beyond our current data availability

Ben 18:31

Were there any milestones and check-ins discussed?

Xijun Zhang 18:34

yes, there will be 5 stages:

S1: Merceros receives data from MF

S2: Merceros initial quality assessment

S3: Cleansing, Ontology applied

S4: Business logic cleansing

S5: Assess quality and load

Xijun Zhang 18:34

so it's kinda a 5x5 matrix for the 5 data domains mentioned

Xijun Zhang 18:34

and each stage will have a milestone

Ben 18:36

Got it, but the plan isn't for them to actually implement or code the copilot them but just to supply the data for it

Xijun Zhang 18:36

it is

Xijun Zhang 18:36

so two projects: 1. copilot, what we discussed yesterday 2. data foundations what I am sharing with you now

Xijun Zhang 18:37

and they have to be married up

Ben 18:37

Ah ok! Thanks for the update!

Xijun Zhang 18:37

I think thanks to the effort you and Steffen are putting in, it's now clear to them the gravitas

Ben 18:38

Do you think we'll be able to keep the 3 week Canon plan?

Ben 18:39

Theyll*

Xijun Zhang 18:40

Michael seems switched on

Xijun Zhang 18:40

let's see if he can manage a team

Ben 18:41

Good to know